

Python for Data Science 2: Guided Machine Learning Project

University: Tek-Up
Group: ING-4-SDIA

Tutor: Haythem Ghazouani

Academic Year 2025-2026

1 Introduction

This 7-week module is designed to immerse students in a professional Machine Learning workflow. Unlike introductory courses, this module focuses on the **end-to-end lifecycle**: from raw data acquisition to a production-ready containerized application.

Module Objective

To master advanced Python for Data Science by implementing a reproducible, scalable, and deployable Machine Learning pipeline integrating MLOps best practices (SMOTE, GridSearchCV, Experiment Tracking, REST APIs).

2 Project Synopsis & Teams Choice

The **Bank Customer Churn** example used in the tutorials serves as a **Project Synopsis** (a comprehensive reference implementation).

Important: Dataset Validation

Students must work in teams. Each team is free to choose their own dataset and objective (e.g., Fraud Detection, Recommendation Systems, Energy Consumption Prediction).

The choice MUST be validated by the tutor before the end of Week 2.

3 Detailed 7-Week Roadmap

3.1 Phase 1: Foundations & Data (Week 1)

- Week 1: Setup, Scraping & EDA:** Environment configuration, Web Scraping (CNBC), and Exploratory Data Analysis.
- Tutorial: tutos/exploration_tuto.pdf*

3.2 Phase 2: The ML Pipeline (Weeks 2-3)

- **Week 2: Preprocessing & Imbalance:** Handling class imbalance (SMOTE) and automated hyperparameter tuning (GridSearchCV).
 - **Week 3: Advanced Modeling & MLflow:** Experiment tracking for XGBoost and Random Forest.
 - *Tutorial: tutos/mlflow_tuto.pdf*
- ## 3.3 Phase 3: Integration & Deployment (Weeks 4-7)
- **Week 4: Backend (FastAPI):** Service monitoring (Health checks) and Batch prediction endpoints.
 - *Tutorial: tutos/fastapi_tuto.pdf*
 - **Week 5: Frontend (React):** Premium dashboard development.
 - *Tutorial: tutos/react_tuto.pdf*
 - **Week 6: Containerization (Docker):** Orchestration using docker compose.
 - *Tutorial: tutos/deployment_tuto.pdf*
 - **Week 7: Final Review & CI/CD (Optional):** Project demonstration.
 - *Tutorial: tutos/cicd_tuto.pdf*

4 Evaluation Criteria

1. **Data Pipeline (20%):** Quality of scraping and EDA insights.
2. **ML Excellence (30%):** Correct use of SMOTE, Pipelines, and MLflow logging.
3. **API UI (30%):** Robustness of the FastAPI endpoints and the React user experience.
4. **Deployment (20%):** Successful execution via Docker.

5 Workspace Structure

Students must strictly adhere to the following directory structure:

- `code/`: Modular Python scripts (`modeling.py`, `app.py`, etc.).
- `frontend/`: React source code and Vite configuration.
- `tutos/`: Lab instructions in PDF format.
- `data/`: Serialized models and dataset samples.
- `Dockerfile.*`: Container definitions.

6 Resources

The reference code, tutorials, and configuration files are available at:
<https://github.com/haythemghz/Python-for-Data-Science-Project>