Python for Data Science 2: Guided Machine Learning Project

Tutor: Haythem Ghazouani

January 28, 2026

1 Introduction

This 7-week module is designed to immerse students in a professional Machine Learning workflow. Unlike introductory courses, this module focuses on the **end-to-end lifecycle**: from raw data acquisition to a production-ready containerized application.

> **Module Objective**
>
> To master advanced Python for Data Science by implementing a reproducible, scalable, and deployable Machine Learning pipeline integrating MLOps best practices (SMOTE, GridSearchCV, Experiment Tracking, REST APIs).

2 Project Synopsis & Teams Choice

The **Bank Customer Churn** example used in the tutorials serves as a **Project Synopsis** (a comprehensive reference implementation).

> **Important: Dataset Validation**
>
> Students must work in teams. Each team is free to choose their own dataset and objective (e.g., Fraud Detection, Recommendation Systems, Energy Consumption Prediction).
> **The choice MUST be validated by the tutor before the end of Week 2.**

3 Detailed 7-Week Roadmap

3.1 Phase 1: Foundations & Data (Week 1)

☐ **Week 1: Setup, Scraping & EDA:** Environment configuration, Web Scraping (CNBC), and Exploratory Data Analysis.

☐ *Tutorial:* `tutos/exploration_tuto.pdf`

3.2 Phase 2: The ML Pipeline (Weeks 2-3)

☐ **Week 2: Preprocessing & Imbalance:** Handling class imbalance (SMOTE) and automated hyperparameter tuning (GridSearchCV).

☐ **Week 3: Advanced Modeling & MLflow:** Experiment tracking for XGBoost and Random Forest.

☐ *Tutorial:* `tutos/mlflow_tuto.pdf`

3.3 Phase 3: Integration & Deployment (Weeks 4-7)

☐ **Week 4: Backend (FastAPI):** Service monitoring (Health checks) and Batch prediction endpoints.

☐ *Tutorial:* `tutos/fastapi_tuto.pdf`

☐ **Week 5: Frontend (React):** Premium dashboard development with responsive state management.

☐ *Tutorial:* `tutos/react_tuto.pdf`

☐ **Week 6: Containerization (Docker):** Orchestration using `docker-compose`.

☐ *Tutorial:* `tutos/deployment_tuto.pdf`

☐ **Week 7: Final Review & CI/CD (Optional):** Project demonstration and automation with GitHub Actions.

☐ *Tutorial:* `tutos/cicd_tuto.pdf`

4 Evaluation Criteria

1. **Data Pipeline (20%):** Quality of scraping and EDA insights.

2. **ML Excellence (30%):** Correct use of SMOTE, Pipelines, and MLflow logging.

3. **API UI (30%):** Robustness of the FastAPI endpoints and the React user experience.

4. **Deployment (20%):** Successful execution via Docker.

5 Workspace Structure

Students must strictly adhere to the following directory structure:

- `code/`: Modular Python scripts (`modeling.py`, `app.py`, etc.).

- `frontend/`: React source code and Vite configuration.

- `tutos/`: Lab instructions in PDF format.

- `data/`: Serialized models and dataset samples.

- `Dockerfile.*`: Container definitions.

6 Resources

The reference code, tutorials, and configuration files are available at:
`https://github.com/haythemghz/Python-for-Data-Science-Project`